

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import sklearn
from sklearn.preprocessing import LabelEncoder
import plotly.express as px
import warnings
warnings.filterwarnings('ignore')

```

```

Hosp=pd.read_csv('Hospitalisation details.csv')
Medic=pd.read_csv('Medical Examinations.csv')
Names=pd.read_excel('Names.xlsx')

```

```
Hosp.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2343 entries, 0 to 2342
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Customer ID           2343 non-null   object
1   year                  2343 non-null   object
2   month                 2343 non-null   object
3   date                  2343 non-null   int64
4   children               2343 non-null   int64
5   charges               2343 non-null   float64
6   Hospital tier          2343 non-null   object
7   City tier              2343 non-null   object
8   State ID              2343 non-null   object
dtypes: float64(1), int64(2), object(6)
memory usage: 164.9+ KB

```

```
Medic.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2335 entries, 0 to 2334
Data columns (total 8 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Customer ID                           2335 non-null   object
1   BMI                                    2335 non-null   float64
2   HBA1C                                  2335 non-null   float64
3   Heart Issues                           2335 non-null   object
4   Any Transplants                        2335 non-null   object
5   Cancer history                         2335 non-null   object
6   NumberOfMajorSurgeries                 2335 non-null   object
7   smoker                                 2335 non-null   object
dtypes: float64(2), object(6)
memory usage: 146.1+ KB

```

```
Names.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2335 entries, 0 to 2334
Data columns (total 2 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Customer ID     2335 non-null   object
1   name            2335 non-null   object
dtypes: object(2)
memory usage: 36.6+ KB
```

Collate the files so that all the information is in one place

```
Customer_details=pd.merge(Hosp,Medic,how='inner',on='Customer ID')
Customer_details
```

	Customer ID	year	month	date	children	charges	Hospital
0	Id2335	1992	Jul	9	0	563.84	tier - 2
1	Id2334	1992	Nov	30	0	570.62	tier - 2
2	Id2333	1993	Jun	30	0	600.00	tier - 2
3	Id2332	1992	Sep	13	0	604.54	tier - 3
4	Id2331	1998	Jul	27	0	637.26	tier - 3
...

2330	Id5	1989	Jun	19	0	55135.40	tier - 1
2331	Id4	1991	Jun	6	1	58571.07	tier - 1
2332	Id3	1970	?	11	3	60021.40	tier - 1
2333	Id2	1977	Jun	8	0	62592.87	tier - 2
2334	Id1	1968	Oct	12	0	63770.43	tier - 1

	City	tier	State	ID	BMI	HBA1C	Heart	Issues	Any	Transplants	\
0	tier - 3		R1013	17.580	4.51		No		No		
1	tier - 1		R1013	17.600	4.39		No		No		
2	tier - 1		R1013	16.470	6.35		No		No		
3	tier - 3		R1013	17.700	6.28		No		No		

4	tier - 3	R1013	22.340	5.57	No	No
...
2330	tier - 2	R1012	35.530	5.45	No	No
2331	tier - 3	R1024	38.095	6.05	No	No
2332	tier - 1	R1012	34.485	11.87	yes	No
2333	tier - 3	R1013	30.360	5.77	No	No
2334	tier - 3	R1013	47.410	7.47	No	No

	Cancer history	NumberOfMajorSurgeries	smoker
0	No		1 No
1	No		1 No
2	Yes		1 No
3	No		1 No
4	No		1 No
...
2330	No	No major surgery	yes
2331	No	No major surgery	yes
2332	No	2	yes
2333	No	No major surgery	yes
2334	No	No major surgery	yes

[2335 rows x 16 columns]

Customer_details.info()

<class 'pandas.core.frame.DataFrame'>

Int64Index: 2335 entries, 0 to 2334

Data columns (total 16 columns):

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	Customer ID	2335 non-null	object
1	year	2335 non-null	object
2	month	2335 non-null	object
3	date	2335 non-null	int64
4	children	2335 non-null	int64
5	charges	2335 non-null	float64
6	Hospital tier	2335 non-null	object
7	City tier	2335 non-null	object
8	State ID	2335 non-null	object
9	BMI	2335 non-null	float64
10	HBA1C	2335 non-null	float64
11	Heart Issues	2335 non-null	object
12	Any Transplants	2335 non-null	object
13	Cancer history	2335 non-null	object
14	NumberOfMajorSurgeries	2335 non-null	object
15	smoker	2335 non-null	object

dtypes: float64(3), int64(2), object(11)

memory usage: 310.1+ KB

```
Customer_details=Customer_details.merge(Names,on='Customer ID')
Customer_details
```

	Customer ID	year	month	date	children	charges	Hospital
tier \							
0	Id2335	1992	Jul	9	0	563.84	tier - 2
1	Id2334	1992	Nov	30	0	570.62	tier - 2
2	Id2333	1993	Jun	30	0	600.00	tier - 2
3	Id2332	1992	Sep	13	0	604.54	tier - 3
4	Id2331	1998	Jul	27	0	637.26	tier - 3
...
2330	Id5	1989	Jun	19	0	55135.40	tier - 1
2331	Id4	1991	Jun	6	1	58571.07	tier - 1
2332	Id3	1970	?	11	3	60021.40	tier - 1
2333	Id2	1977	Jun	8	0	62592.87	tier - 2
2334	Id1	1968	Oct	12	0	63770.43	tier - 1

	City	tier	State	ID	BMI	HBA1C	Heart	Issues	Any	Transplants	\
0	tier - 3		R1013	17.580	4.51		No		No		
1	tier - 1		R1013	17.600	4.39		No		No		
2	tier - 1		R1013	16.470	6.35		No		No		
3	tier - 3		R1013	17.700	6.28		No		No		
4	tier - 3		R1013	22.340	5.57		No		No		
...		
2330	tier - 2		R1012	35.530	5.45		No		No		
2331	tier - 3		R1024	38.095	6.05		No		No		
2332	tier - 1		R1012	34.485	11.87		yes		No		
2333	tier - 3		R1013	30.360	5.77		No		No		
2334	tier - 3		R1013	47.410	7.47		No		No		

	Cancer	history	NumberOfMajorSurgeries	smoker	\
0		No		1	No
1		No		1	No
2		Yes		1	No
3		No		1	No
4		No		1	No
...	
2330		No	No major surgery		yes
2331		No	No major surgery		yes
2332		No		2	yes

2333	No	No major surgery	yes
2334	No	No major surgery	yes

	name
0	German, Mr. Aaron K
1	Rosendahl, Mr. Evan P
2	Albano, Ms. Julie
3	Riveros Gonzalez, Mr. Juan D. Sr.
4	Brietzke, Mr. Jordan
...	...
2330	Kadala, Ms. Kristyn
2331	Osborne, Ms. Kelsey
2332	Lu, Mr. Phil
2333	Lehner, Mr. Matthew D
2334	Hawks, Ms. Kelly

[2335 rows x 17 columns]

```
Customer_details.columns=Customer_details.columns.str.lower()
Customer_details.columns=Customer_details.columns.str.replace(' ','_')
Customer_details.columns
```

```
Index(['customer_id', 'year', 'month', 'date', 'children', 'charges',
       'hospital_tier', 'city_tier', 'state_id', 'bmi', 'hbalc',
       'heart_issues', 'any_transplants', 'cancer_history',
       'numberofmajorsurgeries', 'smoker', 'name'],
      dtype='object')
```

```
Customer_details.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 2335 entries, 0 to 2334
```

```
Data columns (total 17 columns):
```

#	Column	Non-Null Count	Dtype
0	customer_id	2335 non-null	object
1	year	2335 non-null	object
2	month	2335 non-null	object
3	date	2335 non-null	int64
4	children	2335 non-null	int64
5	charges	2335 non-null	float64
6	hospital_tier	2335 non-null	object
7	city_tier	2335 non-null	object
8	state_id	2335 non-null	object
9	bmi	2335 non-null	float64
10	hbalc	2335 non-null	float64
11	heart_issues	2335 non-null	object
12	any_transplants	2335 non-null	object
13	cancer_history	2335 non-null	object
14	numberofmajorsurgeries	2335 non-null	object

```
15 smoker                2335 non-null object
16 name                  2335 non-null object
dtypes: float64(3), int64(2), object(12)
memory usage: 328.4+ KB
```

2. Check for missing values in the dataset

```
Customer_details.isnull().sum()
```

```
customer_id      0
year             0
month            0
date             0
children         0
charges          0
hospital_tier    0
city_tier        0
state_id         0
bmi              0
hbalc            0
heart_issues     0
any_transplants  0
cancer_history   0
numberofmajorsurgeries 0
smoker           0
name             0
dtype: int64
```

The data seems to have trivial values in a few variables. These are "?" in all columns

```
(Customer_details=="?").sum()
```

```
customer_id      0
year             2
month            3
date             0
children         0
charges          0
hospital_tier    1
city_tier        1
state_id         2
bmi              0
hbalc            0
heart_issues     0
any_transplants  0
cancer_history   0
numberofmajorsurgeries 0
smoker           2
```

```

name                                0
dtype: int64

(Customer_details == '?').sum(axis=1).head(20)
0      0
1      0
2      0
3      0
4      0
5      0
6      0
7      0
8      0
9      0
10     0
11     1
12     0
13     1
14     0
15     0
16     0
17     2
18     0
19     0
dtype: int64

Customer_details.shape
(2335, 17)

Customer_details.shape[1]
17

```

3. Find the percentage of rows that have trivial value (for example, ?), and delete such rows if they do not contain significant information

percentage of trivial value in all columns

```

Miss_perc = (Customer_details ==
'?').sum(axis=1) / Customer_details.shape[1] * 100
Miss_perc
Miss_perc[Miss_perc > 0]

```

```

11      5.882353
13      5.882353
17     11.764706
542     5.882353
1046     5.882353
1049     5.882353
1700     5.882353
1775     5.882353
2165     5.882353
2332     5.882353
dtype: float64

Miss_perc[Miss_perc > 0].index

Int64Index([11, 13, 17, 542, 1046, 1049, 1700, 1775, 2165, 2332],
dtype='int64')

```

percentage of trivial value in all rows

```

Miss_perc_rows= (Customer_details==
'?' ).sum(axis=0)/Customer_details.shape[0] * 100
Miss_perc_rows.sort_values(ascending=False)

month      0.128480
state_id   0.085653
year       0.085653
smoker     0.085653
city_tier  0.042827
hospital_tier 0.042827
date       0.000000
children   0.000000
charges    0.000000
name       0.000000
bmi        0.000000
hbalc      0.000000
heart_issues 0.000000
any_transplants 0.000000
cancer_history 0.000000
numberofmajorsurgeries 0.000000
customer_id 0.000000
dtype: float64

Customer_details_noQ=Customer_details.drop(Miss_perc[Miss_perc >
0].index)
Customer_details_noQ

   customer_id  year month  date  children  charges
hospital_tier \
0      Id2335  1992  Jul   9         0    563.84    tier - 2

```


1	Id2334	1992	Nov	30	0	570.62	tier - 2
2	Id2333	1993	Jun	30	0	600.00	tier - 2
3	Id2332	1992	Sep	13	0	604.54	tier - 3
4	Id2331	1998	Jul	27	0	637.26	tier - 3
...
2329	Id6	1962	Aug	4	0	52590.83	tier - 1
2330	Id5	1989	Jun	19	0	55135.40	tier - 1
2331	Id4	1991	Jun	6	1	58571.07	tier - 1
2333	Id2	1977	Jun	8	0	62592.87	tier - 2
2334	Id1	1968	Oct	12	0	63770.43	tier - 1
	city_tier	state_id	bmi	hba1c	heart_issues	any_transplants	\
0	tier - 3	R1013	17.580	4.51	No	No	
1	tier - 1	R1013	17.600	4.39	No	No	
2	tier - 1	R1013	16.470	6.35	No	No	
3	tier - 3	R1013	17.700	6.28	No	No	
4	tier - 3	R1013	22.340	5.57	No	No	
...
2329	tier - 3	R1011	32.800	6.59	No	No	
2330	tier - 2	R1012	35.530	5.45	No	No	
2331	tier - 3	R1024	38.095	6.05	No	No	
2333	tier - 3	R1013	30.360	5.77	No	No	
2334	tier - 3	R1013	47.410	7.47	No	No	
	cancer_history	numberofmajorsurgeries	smoker	\			
0	No		1	No			
1	No		1	No			
2	Yes		1	No			
3	No		1	No			
4	No		1	No			
...			
2329	No	No major surgery	yes				
2330	No	No major surgery	yes				
2331	No	No major surgery	yes				
2333	No	No major surgery	yes				
2334	No	No major surgery	yes				
	name						
0	German, Mr. Aaron K						
1	Rosendahl, Mr. Evan P						

```

2           Albano, Ms. Julie
3  Riveros Gonzalez, Mr. Juan D. Sr.
4           Brietzke, Mr. Jordan
...
2329        Baker, Mr. Russell B.
2330        Kadala, Ms. Kristyn
2331        Osborne, Ms. Kelsey
2333        Lehner, Mr. Matthew D
2334        Hawks, Ms. Kelly

[2325 rows x 17 columns]

Customer_details_noQ.columns

Index(['customer_id', 'year', 'month', 'date', 'children', 'charges',
      'hospital_tier', 'city_tier', 'state_id', 'bmi', 'hbalc',
      'heart_issues', 'any_transplants', 'cancer_history',
      'numberofmajorsurgeries', 'smoker', 'name'],
      dtype='object')

```

4. Use the necessary transformation methods to deal with the nominal and ordinal categorical variables in the dataset

Label Encoding

Nominal Variable - State ID

Ordinal Variable - Ranking variables-numbers are assigned to categories based on rank - 'Hospital tier', 'City tier'

```

Customer_details_noQ['hospital_tier'].unique()

array(['tier - 2', 'tier - 3', 'tier - 1'], dtype=object)

Customer_details_noQ.groupby('hospital_tier').count()['customer_id']

hospital_tier
tier - 1      300
tier - 2     1334
tier - 3      691
Name: customer_id, dtype: int64

import sklearn
from sklearn.preprocessing import LabelEncoder

```

```

label_encoder=LabelEncoder()
Customer_details_noQ['hospital_tier_ord']=label_encoder.fit_transform(
Customer_details_noQ['hospital_tier'])

Customer_details_noQ['hospital_tier_ord'].unique()

array([1, 2, 0])

Customer_details_noQ['hospital_tier_ord']
0      1
1      1
2      1
3      2
4      2
..
2329   0
2330   0
2331   0
2333   1
2334   0
Name: hospital_tier_ord, Length: 2325, dtype: int32

pd.crosstab(Customer_details_noQ['hospital_tier_ord'],Customer_details
_noQ['hospital_tier'])

hospital_tier      tier - 1  tier - 2  tier - 3
hospital_tier_ord
0                300         0         0
1                 0       1334         0
2                 0         0       691

Customer_details_noQ['city_tier'].unique()

array(['tier - 3', 'tier - 1', 'tier - 2'], dtype=object)

Customer_details_noQ.groupby('city_tier').count()['customer_id']

city_tier
tier - 1      729
tier - 2      807
tier - 3      789
Name: customer_id, dtype: int64

Customer_details_noQ['city_tier_ord']=
label_encoder.fit_transform(Customer_details_noQ['city_tier'])
Customer_details_noQ['city_tier_ord']

0      2
1      0
2      0
3      2

```

```

4      2
2329   2
2330   1
2331   2
2333   2
2334   2
Name: city_tier_ord, Length: 2325, dtype: int32

Customer_details_noQ['city_tier_ord'].unique()

array([2, 0, 1])

pd.crosstab(Customer_details_noQ['city_tier_ord'],Customer_details_noQ
['city_tier'])

city_tier      tier - 1  tier - 2  tier - 3
city_tier_ord
0              729         0         0
1               0        807         0
2               0         0        789

Customer_details_noQ.head(5)

   customer_id  year month  date  children  charges  hospital_tier
city_tier \
0      Id2335  1992   Jul    9         0   563.84      tier - 2  tier
- 3
1      Id2334  1992  Nov   30         0   570.62      tier - 2  tier
- 1
2      Id2333  1993   Jun   30         0   600.00      tier - 2  tier
- 1
3      Id2332  1992   Sep   13         0   604.54      tier - 3  tier
- 3
4      Id2331  1998   Jul   27         0   637.26      tier - 3  tier
- 3

   state_id  bmi  hba1c  heart_issues  any_transplants
cancer_history \
0    R1013  17.58   4.51             No              No              No
1    R1013  17.60   4.39             No              No              No
2    R1013  16.47   6.35             No              No              Yes
3    R1013  17.70   6.28             No              No              No
4    R1013  22.34   5.57             No              No              No

   numberofmajorsurgeries  smoker  name \

```

0	1	No	German, Mr.	Aaron K
1	1	No	Rosendahl, Mr.	Evan P
2	1	No	Albano, Ms.	Julie
3	1	No	Riveros Gonzalez, Mr.	Juan D. Sr.
4	1	No	Brietzke, Mr.	Jordan

	hospital_tier_ord	city_tier_ord
0	1	2
1	1	0
2	1	0
3	2	2
4	2	2

5. The dataset has State ID, which has around 16 states. All states are not represented in equal proportions in the data. Creating dummy variables for all regions may also result in too many insignificant predictors. Nevertheless, only R1011, R1012, and R1013 are worth investigating further. Design a suitable strategy to create dummy variables with these restraints.

Creating dummy variable of all 16 states will may lead to insignificant predictors.

####Choosing the most frequent of each category

```
Customer_details_noQ['state_id'].unique()

array(['R1013', 'R1012', 'R1011', 'R1015', 'R1019', 'R1016', 'R1018',
       'R1025', 'R1024', 'R1023', 'R1014', 'R1021', 'R1017', 'R1020',
       'R1026', 'R1022'], dtype=object)

SC=Customer_details_noQ.groupby('state_id').count()
['customer_id'].sort_values(ascending=False)
SC
```

```

state_id
R1013    609
R1011    574
R1012    572
R1024    159
R1026     84
R1021     70
R1016     64
R1025     40
R1023     38
R1017     36
R1019     26
R1022     14
R1014     13
R1015     11
R1018      9
R1020      6
Name: customer_id, dtype: int64

SC[:3].index    # first 3 state are most frequent

Index(['R1013', 'R1011', 'R1012'], dtype='object', name='state_id')

for i in SC[:3].index:
    var_name= 'State_ID_'+ i
    print(var_name)
    Customer_details_noQ[var_name]=0
    Customer_details_noQ.loc[Customer_details_noQ['state_id'] ==
i,var_name]=1

State_ID_R1013
State_ID_R1011
State_ID_R1012

#Customer_details_noQ['State ID']
Customer_details_noQ['State_ID_R1013'].value_counts()

0    1716
1     609
Name: State_ID_R1013, dtype: int64

Customer_details_noQ['State_ID_R1012'].value_counts()

0    1753
1     572
Name: State_ID_R1012, dtype: int64

Customer_details_noQ['State_ID_R1011'].value_counts()

```

```
0    1751
1     574
Name: State_ID_R1011, dtype: int64
```

6. The variable NumberOfMajorSurgeries also appears to have string values. Apply a suitable method to clean up this variable.

```
Customer_details_noQ['numberofmajorsurgeries'].unique()
array(['1', 'No major surgery', '2', '3'], dtype=object)

Customer_details_noQ['numberofmajorsurgeries'].value_counts()
No major surgery    1070
1                   961
2                   272
3                    22
Name: numberofmajorsurgeries, dtype: int64

Customer_details_noQ['numberofmajorsurgeries']=pd.to_numeric(Customer_
details_noQ['numberofmajorsurgeries'],errors='coerce')
Customer_details_noQ['numberofmajorsurgeries']

0         1.0
1         1.0
2         1.0
3         1.0
4         1.0
...
2329      NaN
2330      NaN
2331      NaN
2333      NaN
2334      NaN
Name: numberofmajorsurgeries, Length: 2325, dtype: float64

Customer_details_noQ['numberofmajorsurgeries'].fillna(0,inplace=True)
Customer_details_noQ['numberofmajorsurgeries'].unique()
array([1., 0., 2., 3.]
```

7. Age appears to be a significant factor in this analysis. Calculate the patients' ages based on their

dates of birth.

```
Customer_details_noQ.year=Customer_details_noQ.year.astype(int)
Customer_details_noQ['age']=2024-Customer_details_noQ.year
Customer_details_noQ['age']
```

```
0      32
1      32
2      31
3      32
4      26
```

```
..
```

```
2329   62
2330   35
2331   33
2333   47
2334   56
```

```
Name: age, Length: 2325, dtype: int32
```

8. The gender of the patient may be an important factor in determining the cost of hospitalization. The salutations in a beneficiary's name can be used to determine their gender. Make a new field for the beneficiary's gender.

```
Customer_details_noQ['title']=Customer_details_noQ['name'].str.split(' ',1)[1].str.strip()
Customer_details_noQ['title'].value_counts()
```

```
Mr      1160
Ms       1023
Mrs       142
```

```
Name: title, dtype: int64
```

```
Customer_details_noQ.shape
```



```
(2325, 24)
```

```
Customer_details_noQ [ 'gender'] = 'Female'  
Customer_details_noQ.loc[Customer_details_noQ.title == 'Mr' ,  
'gender'] = 'Male'
```

```
Customer_details_noQ['gender']
```

```
0      Male  
1      Male  
2     Female  
3      Male  
4      Male
```

```
...  
2329     Male  
2330    Female  
2331    Female  
2333     Male  
2334    Female
```

```
Name: gender, Length: 2325, dtype: object
```

```
Customer_details_noQ.columns
```

```
Index(['customer_id', 'year', 'month', 'date', 'children', 'charges',  
      'hospital_tier', 'city_tier', 'state_id', 'bmi', 'hbalc',  
      'heart_issues', 'any_transplants', 'cancer_history',  
      'numberofmajorsurgeries', 'smoker', 'name',  
      'hospital_tier_ord',  
      'city_tier_ord', 'State_ID_R1013', 'State_ID_R1011',  
      'State_ID_R1012',  
      'age', 'title', 'gender'],  
      dtype='object')
```

```
Customer_details_noQ
```

	customer_id	year	month	date	children	charges	
hospital_tier \							
0	Id2335	1992	Jul	9	0	563.84	tier - 2
1	Id2334	1992	Nov	30	0	570.62	tier - 2
2	Id2333	1993	Jun	30	0	600.00	tier - 2
3	Id2332	1992	Sep	13	0	604.54	tier - 3
4	Id2331	1998	Jul	27	0	637.26	tier - 3
...
2329	Id6	1962	Aug	4	0	52590.83	tier - 1

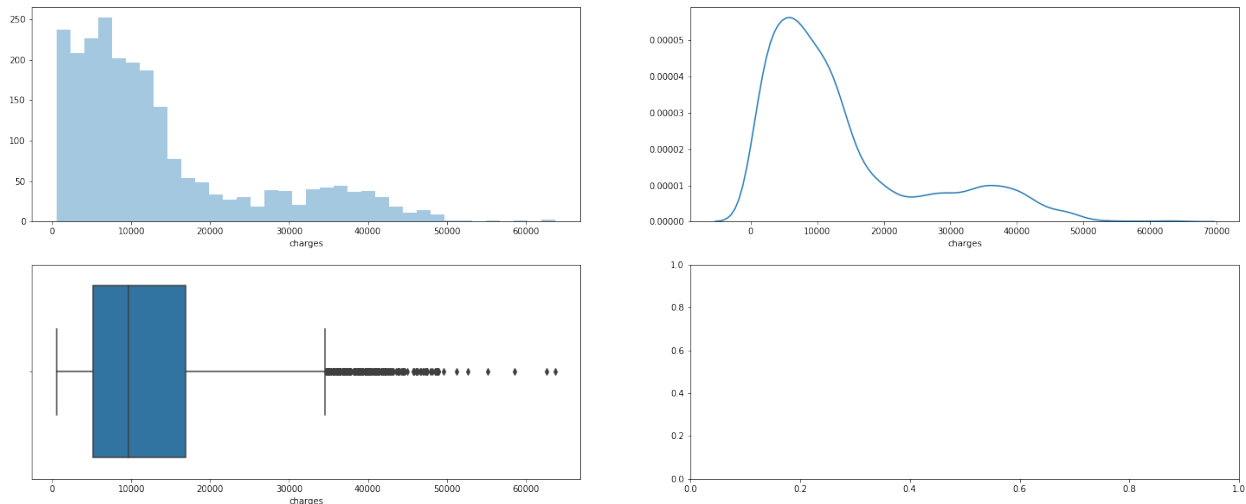
2330	Id5	1989	Jun	19	0	55135.40	tier - 1
2331	Id4	1991	Jun	6	1	58571.07	tier - 1
2333	Id2	1977	Jun	8	0	62592.87	tier - 2
2334	Id1	1968	Oct	12	0	63770.43	tier - 1
	city_tier	state_id	bmi	...	smoker	\	
0	tier - 3	R1013	17.580	...	No		
1	tier - 1	R1013	17.600	...	No		
2	tier - 1	R1013	16.470	...	No		
3	tier - 3	R1013	17.700	...	No		
4	tier - 3	R1013	22.340	...	No		
...		
2329	tier - 3	R1011	32.800	...	yes		
2330	tier - 2	R1012	35.530	...	yes		
2331	tier - 3	R1024	38.095	...	yes		
2333	tier - 3	R1013	30.360	...	yes		
2334	tier - 3	R1013	47.410	...	yes		
					name	hospital_tier_ord	
city_tier_ord	\						
0	German, Mr. Aaron K				1		
2							
1	Rosendahl, Mr. Evan P				1		
0							
2	Albano, Ms. Julie				1		
0							
3	Riveros Gonzalez, Mr. Juan D. Sr.				2		
2							
4	Brietzke, Mr. Jordan				2		
2							
...		
..							
2329	Baker, Mr. Russell B.				0		
2							
2330	Kadala, Ms. Kristyn				0		
1							
2331	Osborne, Ms. Kelsey				0		
2							
2333	Lehner, Mr. Matthew D				1		
2							
2334	Hawks, Ms. Kelly				0		
2							
	State_ID_R1013	State_ID_R1011	State_ID_R1012	age	title	gender	
0	1	0	0	32	Mr	Male	

1	1	0	0	32	Mr	Male
2	1	0	0	31	Ms	Female
3	1	0	0	32	Mr	Male
4	1	0	0	26	Mr	Male
...
2329	0	1	0	62	Mr	Male
2330	0	0	1	35	Ms	Female
2331	0	0	0	33	Ms	Female
2333	1	0	0	47	Mr	Male
2334	1	0	0	56	Ms	Female

[2325 rows x 25 columns]

9. You should also visualize the distribution of costs using a histogram, box and whisker plot, and swarm plot.


```
fig,ax=plt.subplots(2,2,figsize=[25,10])
sns.distplot(Customer_details_noQ['charges'],hist=True,kde=
False,ax=ax[0][0])
sns.distplot(Customer_details_noQ['charges'],hist=False,kde=
True,ax=ax[0][1])
sns.boxplot(Customer_details_noQ['charges'],ax=ax[1][0])
plt.show()
```



10. State how the distribution is different across gender and tiers of hospitals


WRT Gender

```
plt.figure(figsize=(15,5))
sns.boxplot(x='charges',y='gender',data=Customer_details_noQ )
plt.show()
```




WRT city tier

```
plt.figure(figsize=(15,5))
sns.boxplot(x='charges',y='city_tier',data=Customer_details_noQ )
plt.show()
```



WRT Hospital tier

```
plt.figure(figsize=(15,5))
sns.boxplot(x='charges',y='hospital_tier',data=Customer_details_noQ )
plt.show()
```



11. Create a radar chart to showcase the median hospitalization cost for each tier of hospitals

```
median = Customer_details_noQ.groupby('hospital_tier')
[['charges']].median().reset_index()
median
```

	hospital_tier	charges
0	tier - 1	32097.435
1	tier - 2	7168.760
2	tier - 3	10676.830

```

fig = px.line_polar(median, r='charges', theta='hospital_tier') #,
line_close=True
fig.show()

{"config":{"plotlyServerURL":"https://plot.ly"},"data":
[{"hovertemplate":"charges=%{r}<br>hospital_tier=%{theta}<extra></
extra>","legendgroup":"","line":
{"color":"#636efa","dash":"solid"},"marker":
{"symbol":"circle"},"mode":"lines","name":"","r":
[32097.434999999998,7168.76,10676.83],"showlegend":false,"subplot":"po
lar","theta":["tier - 1","tier - 2","tier -
3"],"type":"scatterpolar"}],"layout":{"legend":
{"tracegroupgap":0},"margin":{"t":60},"polar":{"angularaxis":
{"direction":"clockwise","rotation":90},"domain":{"x":[0,1],"y":
[0,1]},"template":{"data":{"bar":{"error_x":
{"color":"#2a3f5f"},"error_y":{"color":"#2a3f5f"},"marker":{"line":
{"color":"#E5ECF6","width":0.5},"pattern":
{"fillmode":"overlay","size":10,"solidity":0.2},"type":"bar"}],"barpo
lar":[{"marker":{"line":{"color":"#E5ECF6","width":0.5},"pattern":
{"fillmode":"overlay","size":10,"solidity":0.2},"type":"barpolar"}],"
carpet":[{"aaxis":
{"endlinecolor":"#2a3f5f","gridcolor":"white","linecolor":"white","min
orgridcolor":"white","startlinecolor":"#2a3f5f"},"baxis":
{"endlinecolor":"#2a3f5f","gridcolor":"white","linecolor":"white","min
orgridcolor":"white","startlinecolor":"#2a3f5f"},"type":"carpet"}],"ch
oropleth":[{"colorbar":
{"outlinewidth":0,"ticks":"","type":"choropleth"}],"contour":
[{"colorbar":{"outlinewidth":0,"ticks":"","colorscale":
[[0,"#0d0887"],[0.11111111111111111,"#46039f"],
[0.22222222222222222,"#7201a8"],[0.33333333333333333,"#9c179e"],
[0.44444444444444444,"#bd3786"],[0.55555555555555556,"#d8576b"],
[0.66666666666666666,"#ed7953"],[0.77777777777777778,"#fb9f3a"],
[0.88888888888888888,"#fdca26"],
[1,"#f0f921"]],"type":"contour"}],"contourcarpet":[{"colorbar":
{"outlinewidth":0,"ticks":"","type":"contourcarpet"}],"heatmap":
[{"colorbar":{"outlinewidth":0,"ticks":"","colorscale":
[[0,"#0d0887"],[0.11111111111111111,"#46039f"],
[0.22222222222222222,"#7201a8"],[0.33333333333333333,"#9c179e"],
[0.44444444444444444,"#bd3786"],[0.55555555555555556,"#d8576b"],
[0.66666666666666666,"#ed7953"],[0.77777777777777778,"#fb9f3a"],
[0.88888888888888888,"#fdca26"],
[1,"#f0f921"]],"type":"heatmap"}],"heatmapgl":[{"colorbar":
{"outlinewidth":0,"ticks":"","colorscale":[[0,"#0d0887"],
[0.11111111111111111,"#46039f"],[0.22222222222222222,"#7201a8"],
[0.33333333333333333,"#9c179e"],[0.44444444444444444,"#bd3786"],
[0.55555555555555556,"#d8576b"],[0.66666666666666666,"#ed7953"],
[0.77777777777777778,"#fb9f3a"],[0.88888888888888888,"#fdca26"],
[1,"#f0f921"]],"type":"heatmapgl"}],"histogram":[{"marker":{"pattern":
{"fillmode":"overlay","size":10,"solidity":0.2},"type":"histogram"}],
"histogram2d":[{"colorbar":{"outlinewidth":0,"ticks":"","colorscale":

```

```

[[0,"#0d0887"],[0.1111111111111111,"#46039f"],
[0.2222222222222222,"#7201a8"],[0.3333333333333333,"#9c179e"],
[0.4444444444444444,"#bd3786"],[0.5555555555555556,"#d8576b"],
[0.6666666666666666,"#ed7953"],[0.7777777777777778,"#fb9f3a"],
[0.8888888888888888,"#fdca26"],
[1,"#f0f921"]], "type": "histogram2d"}], "histogram2dcontour":
[{"colorbar": {"linewidth": 0, "ticks": ""}, "colorscale":
[[0,"#0d0887"],[0.1111111111111111,"#46039f"],
[0.2222222222222222,"#7201a8"],[0.3333333333333333,"#9c179e"],
[0.4444444444444444,"#bd3786"],[0.5555555555555556,"#d8576b"],
[0.6666666666666666,"#ed7953"],[0.7777777777777778,"#fb9f3a"],
[0.8888888888888888,"#fdca26"],
[1,"#f0f921"]], "type": "histogram2dcontour"}], "mesh3d": [{"colorbar":
{"linewidth": 0, "ticks": ""}, "type": "mesh3d"}], "parcoords": [{"line":
{"colorbar": {"linewidth": 0, "ticks": ""}, "type": "parcoords"}], "pie":
[{"automargin": true, "type": "pie"}], "scatter": [{"fillpattern":
{"fillmode": "overlay", "size": 10, "solidity": 0.2}, "type": "scatter"}], "scatter3d": [{"line": {"colorbar": {"linewidth": 0, "ticks": ""}, "marker":
{"colorbar":
{"linewidth": 0, "ticks": ""}, "type": "scatter3d"}], "scattercarpet":
[{"marker": {"colorbar":
{"linewidth": 0, "ticks": ""}, "type": "scattercarpet"}], "scattergeo":
[{"marker": {"colorbar":
{"linewidth": 0, "ticks": ""}, "type": "scattergeo"}], "scattergl":
[{"marker": {"colorbar":
{"linewidth": 0, "ticks": ""}, "type": "scattergl"}], "scattermapbox":
[{"marker": {"colorbar":
{"linewidth": 0, "ticks": ""}, "type": "scattermapbox"}], "scatterpolar":
[{"marker": {"colorbar":
{"linewidth": 0, "ticks": ""}, "type": "scatterpolar"}], "scatterpolargl":
[{"marker": {"colorbar":
{"linewidth": 0, "ticks": ""}, "type": "scatterpolargl"}], "scatterternary":
[{"marker": {"colorbar":
{"linewidth": 0, "ticks": ""}, "type": "scatterternary"}], "surface":
[{"colorbar": {"linewidth": 0, "ticks": ""}, "colorscale":
[[0,"#0d0887"],[0.1111111111111111,"#46039f"],
[0.2222222222222222,"#7201a8"],[0.3333333333333333,"#9c179e"],
[0.4444444444444444,"#bd3786"],[0.5555555555555556,"#d8576b"],
[0.6666666666666666,"#ed7953"],[0.7777777777777778,"#fb9f3a"],
[0.8888888888888888,"#fdca26"],
[1,"#f0f921"]], "type": "surface"}], "table": [{"cells": {"fill":
{"color": "#EBF0F8"}, "line": {"color": "white"}}, "header": {"fill":
{"color": "#C8D4E3"}, "line":
{"color": "white"}}, "type": "table"}], "layout": {"annotationdefaults":
{"arrowcolor": "#2a3f5f", "arrowhead": 0, "arrowwidth": 1}, "autotypenumbers": "strict", "coloraxis": {"colorbar":
{"linewidth": 0, "ticks": ""}, "colorscale": {"diverging":
[[0,"#8e0152"],[0.1,"#c51b7d"],[0.2,"#de77ae"],[0.3,"#f1b6da"],
[0.4,"#fde0ef"],[0.5,"#f7f7f7"],[0.6,"#e6f5d0"],[0.7,"#b8e186"],
[0.8,"#7fbcb4"],[0.9,"#4d9221"],[1,"#276419"]], "sequential":

```

```

[[0,"#0d0887"],[0.1111111111111111,"#46039f"],
[0.2222222222222222,"#7201a8"],[0.3333333333333333,"#9c179e"],
[0.4444444444444444,"#bd3786"],[0.5555555555555556,"#d8576b"],
[0.6666666666666666,"#ed7953"],[0.7777777777777778,"#fb9f3a"],
[0.8888888888888888,"#fdca26"],[1,"#f0f921"]],"sequentialminus":
[[0,"#0d0887"],[0.1111111111111111,"#46039f"],
[0.2222222222222222,"#7201a8"],[0.3333333333333333,"#9c179e"],
[0.4444444444444444,"#bd3786"],[0.5555555555555556,"#d8576b"],
[0.6666666666666666,"#ed7953"],[0.7777777777777778,"#fb9f3a"],
[0.8888888888888888,"#fdca26"],[1,"#f0f921"]]],"colorway":
["#636efa","#EF553B","#00cc96","#ab63fa","#FFA15A","#19d3f3","#FF6692",
"#B6E880","#FF97FF","#FECB52"],"font":{"color":"#2a3f5f"},"geo":
{"bgcolor":"white","lakecolor":"white","landcolor":"#E5ECF6","showlake
s":true,"showland":true,"subunitcolor":"white"},"hoverlabel":
{"align":"left"},"hovermode":"closest","mapbox":
{"style":"light"},"paper_bgcolor":"white","plot_bgcolor":"#E5ECF6","po
lar":{"angularaxis":
{"gridcolor":"white","linecolor":"white","ticks":""},"bgcolor":"#E5ECF
6","radialaxis":
{"gridcolor":"white","linecolor":"white","ticks":""}}},"scene":
{"xaxis":
{"backgroundcolor":"#E5ECF6","gridcolor":"white","gridwidth":2,"lineco
lor":"white","showbackground":true,"ticks":"","zerolinecolor":"white"}
,"yaxis":
{"backgroundcolor":"#E5ECF6","gridcolor":"white","gridwidth":2,"lineco
lor":"white","showbackground":true,"ticks":"","zerolinecolor":"white"}
,"zaxis":
{"backgroundcolor":"#E5ECF6","gridcolor":"white","gridwidth":2,"lineco
lor":"white","showbackground":true,"ticks":"","zerolinecolor":"white"}
},"shapedefaults":{"line":{"color":"#2a3f5f"}},ternary":{"aaxis":
{"gridcolor":"white","linecolor":"white","ticks":""},"baxis":
{"gridcolor":"white","linecolor":"white","ticks":""},"bgcolor":"#E5ECF
6","caxis":
{"gridcolor":"white","linecolor":"white","ticks":""}}},"title":
{"x":5.0e-2},"xaxis":
{"automargin":true,"gridcolor":"white","linecolor":"white","ticks":"","
"title":
{"standoff":15},"zerolinecolor":"white","zerolinewidth":2},"yaxis":
{"automargin":true,"gridcolor":"white","linecolor":"white","ticks":"","
"title":{"standoff":15},"zerolinecolor":"white","zerolinewidth":2}}}}}

```


12. Create a frequency table and a stacked bar chart to visualize the count of people in the different tiers of cities and hospitals

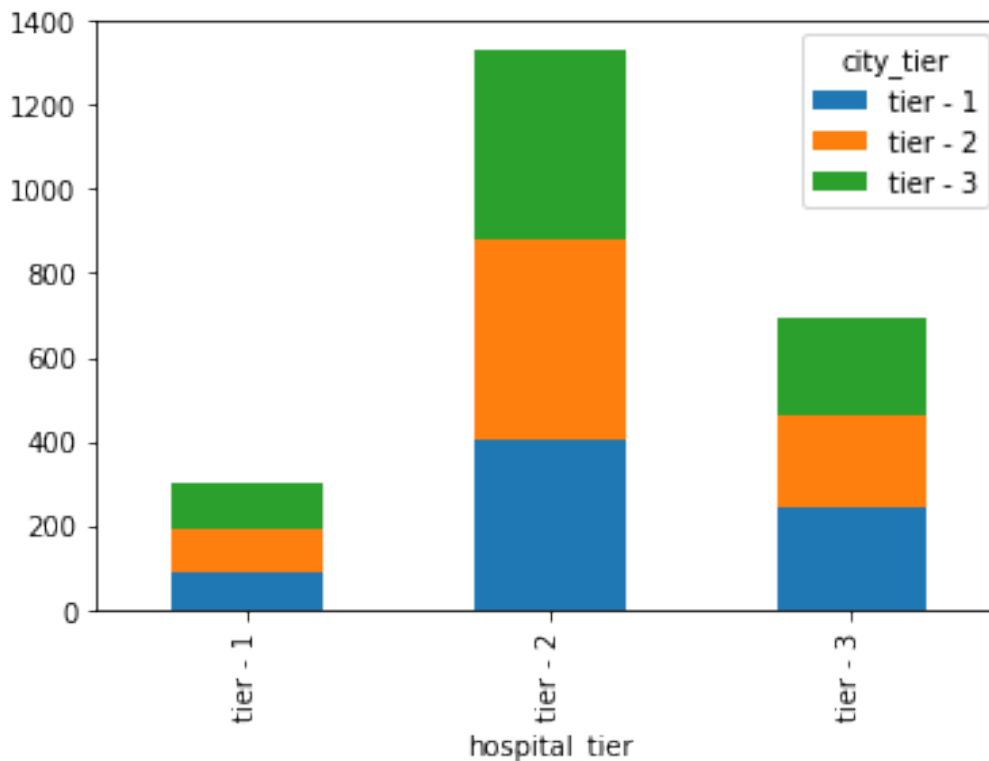
```
pd.crosstab(Customer_details_noQ['hospital_tier'], Customer_details_noQ['city_tier'])
```

city_tier	tier - 1	tier - 2	tier - 3
hospital_tier			
tier - 1	85	106	109
tier - 2	403	479	452
tier - 3	241	222	228

```
#plt.figure(figsize=[12,6])
```

```
pd.crosstab(Customer_details_noQ['hospital_tier'], Customer_details_noQ['city_tier']).plot.bar(stacked=True)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x17cfa20efc8>
```



13. Test the following null hypotheses:

- #a. The average hospitalization costs for the three types of hospitals are not significantly different.
- #b. The average hospitalization costs for the three types of cities are not significantly different.
- #c. The average hospitalization cost for smokers is not significantly different from the average cost for nonsmokers.
- #d. Smoking and heart issues are independent.

a. The average hospitalization costs for the three types of hospitals are not significantly different.

```
from scipy.stats import ttest_1samp

Customer_details_noQ.columns

Index(['customer_id', 'year', 'month', 'date', 'children', 'charges',
      'hospital_tier', 'city_tier', 'state_id', 'bmi', 'hbalc',
      'heart_issues', 'any_transplants', 'cancer_history',
      'numberofmajorsurgeries', 'smoker', 'name',
      'hospital_tier_ord',
      'city_tier_ord', 'State_ID_R1013', 'State_ID_R1011',
      'State_ID_R1012',
      'age', 'title', 'gender'],
      dtype='object')
```

Anova Test

```
import statsmodels.api as sm
from statsmodels.formula.api import ols

model = ols('charges ~ hospital_tier', data=Customer_details_noQ).fit()
res = sm.stats.anova_lm(model)
res
```

	df	sum_sq	mean_sq	F
PR(>F)				
hospital_tier	2.0	9.763011e+10	4.881505e+10	493.989566
1.773822e-179				
Residual	2322.0	2.294554e+11	9.881799e+07	NaN
NaN				

looking into `pvalue(1.77) > alpha(0.005)` , we cannot reject null hypothesis and can conclude costs for the three types of hospitals are significantly different.

```
File "<ipython-input-65-2e384f72ed92>", line 1
    looking into pvalue(1.77) > alpha(0.005) , we cannot reject null
hypothesis and
              ^
SyntaxError: invalid syntax
```

b. The average hospitalization costs for the three types of cities are not significantly different.

```
model = ols('charges ~ city_tier',data=Customer_details_noQ).fit()
res = sm.stats.anova_lm(model)
res
```

looking into `pvalue(0.233) > alpha(0.005)` , we cannot reject null hypothesis and can conclude costs for the three types of cities are significantly different.

c. The average hospitalization cost for smokers is not significantly different from the average cost for nonsmokers.

T-test

```
cost_for_smokers=Customer_details_noQ.loc[Customer_details_noQ.smoker
== 'yes','charges']
cost_for_nonsmokers=Customer_details_noQ.loc[Customer_details_noQ.smok
er != 'yes','charges']
print(cost_for_smokers.count())
print(cost_for_nonsmokers.count())

import numpy as np
from scipy.stats import ttest_ind
```

```
Stat,P_value=ttest_ind(cost_for_smokers,cost_for_nonsmokers)
print('P value to check hypothesis ',P_value)
```

Looking at the $p_value < \alpha(0.05)$, we can reject the null hypothesis and conclude that Average hospitalization cost for smokers is significantly different than non-smokers

d. Smoking and heart issues are independent.

Chi-squared test

```
observed =pd.crosstab(Customer_details_noQ.smoker,
Customer_details_noQ.heart_issues)
observed

from scipy.stats import chi2_contingency

# Defining the observed frequencies
#observed = [Customer_details_noQ.smoker,
Customer_details_noQ.heart_issues] string will not be considered

# Perform chi-squared test
chi2, p, dof, expected = chi2_contingency(observed)

print("Observed frequencies:")
print(observed)
print("Expected frequencies:")
print(expected)
print("Chi-squared statistic:", chi2)
print("p-value:", p)
```

#observation of Chi Test: Pvalue(0.76) > alpha(0.05) cannot reject Null Hypothesis. means smoking and heart issue are independent event

----- * END OF
FIRST PART OF PROJECT * -----

Machine learning

```
Customer_details_noQ.columns
```

```
Index(['customer_id', 'year', 'month', 'date', 'children', 'charges',
      'hospital_tier', 'city_tier', 'state_id', 'bmi', 'hbalc',
      'heart_issues', 'any_transplants', 'cancer_history',
      'numberofmajorsurgeries', 'smoker', 'name',
      'hospital_tier_ord',
      'city_tier_ord', 'State_ID_R1013', 'State_ID_R1011',
      'State_ID_R1012',
      'age', 'title', 'gender'],
      dtype='object')
```

```
Customer_details_noQ.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 2325 entries, 0 to 2334
```

```
Data columns (total 25 columns):
```

#	Column	Non-Null Count	Dtype
0	customer_id	2325 non-null	object
1	year	2325 non-null	int32
2	month	2325 non-null	object
3	date	2325 non-null	int64
4	children	2325 non-null	int64
5	charges	2325 non-null	float64
6	hospital_tier	2325 non-null	object
7	city_tier	2325 non-null	object
8	state_id	2325 non-null	object
9	bmi	2325 non-null	float64
10	hbalc	2325 non-null	float64
11	heart_issues	2325 non-null	object
12	any_transplants	2325 non-null	object
13	cancer_history	2325 non-null	object
14	numberofmajorsurgeries	2325 non-null	float64
15	smoker	2325 non-null	object
16	name	2325 non-null	object
17	hospital_tier_ord	2325 non-null	int32
18	city_tier_ord	2325 non-null	int32
19	State_ID_R1013	2325 non-null	int64
20	State_ID_R1011	2325 non-null	int64
21	State_ID_R1012	2325 non-null	int64
22	age	2325 non-null	int32
23	title	2325 non-null	object
24	gender	2325 non-null	object

```
dtypes: float64(4), int32(4), int64(5), object(12)
```

```
memory usage: 515.9+ KB
```

Problem - 1

- Examine the correlation between predictors to identify highly correlated predictors

```
data = Customer_details_noQ[[ 'children', 'charges',
                              'bmi', 'hbalc', 'numberofmajorsurgeries', 'hospital_tier_ord',
                              'city_tier_ord', 'State_ID_R1013', 'State_ID_R1011',
                              'State_ID_R1012',
                              'age', ]]

```

```
data_corr=data.corr()
data_corr
```

	children	charges	bmi	hbalc	\
children	1.000000	0.055901	-0.005339	-0.101379	
charges	0.055901	1.000000	0.346730	0.139697	
bmi	-0.005339	0.346730	1.000000	-0.006920	
hbalc	-0.101379	0.139697	-0.006920	1.000000	
numberofmajorsurgeries	-0.113161	0.053308	0.018851	-0.091594	
hospital_tier_ord	-0.052438	-0.446687	-0.104771	0.057855	
city_tier_ord	-0.015760	0.035300	0.038123	-0.005404	
State_ID_R1013	-0.013834	-0.150634	-0.208744	0.033453	
State_ID_R1011	0.011666	0.286956	0.115671	0.015525	
State_ID_R1012	0.005247	-0.074636	0.017939	-0.019513	
age	-0.005457	0.304395	0.049260	0.460558	


	numberofmajorsurgeries	hospital_tier_ord	\
children	-0.113161	-0.052438	
charges	0.053308	-0.446687	
bmi	0.018851	-0.104771	
hbalc	-0.091594	0.057855	
numberofmajorsurgeries	1.000000	0.033230	
hospital_tier_ord	0.033230	1.000000	
city_tier_ord	0.027937	-0.039755	
State_ID_R1013	-0.002056	0.002455	
State_ID_R1011	0.000208	-0.114685	
State_ID_R1012	-0.002098	0.020272	
age	0.151442	0.133771	

	city_tier_ord	State_ID_R1013	State_ID_R1011
children	-0.015760	-0.013834	0.011666
charges	0.035300	-0.150634	0.286956
bmi	0.038123	-0.208744	0.115671
hbalc	-0.005404	0.033453	0.015525
numberofmajorsurgeries	0.027937	-0.002056	0.000208
hospital_tier_ord	-0.039755	0.002455	-0.114685
city_tier_ord	1.000000	0.002766	0.036049

State_ID_R1013	0.002766	1.000000	-0.341085
State_ID_R1011	0.036049	-0.341085	1.000000
State_ID_R1012	-0.018253	-0.340296	-0.327054
age	-0.008070	-0.011926	0.008022

	State_ID_R1012	age
children	0.005247	-0.005457
charges	-0.074636	0.304395
bmi	0.017939	0.049260
hba1c	-0.019513	0.460558
numberofmajorsurgeries	-0.002098	0.151442
hospital_tier_ord	0.020272	0.133771
city_tier_ord	-0.018253	-0.008070
State_ID_R1013	-0.340296	-0.011926
State_ID_R1011	-0.327054	0.008022
State_ID_R1012	1.000000	-0.005229
age	-0.005229	1.000000

```
plt.figure(figsize=(15,5))
ma = np.ones_like(data_corr)
ma[np.tril_indices_from(ma)]=0
sns.heatmap(data_corr,annot=True,mask=ma,cmap='PuRd')
plt.xticks(rotation=30)
plt.show()
```



Problem - 2

Develop a regression model Linear or Ridge. Evaluate the model with k-fold cross validation. Also, ensure that you apply all the following suggestions:

- Implement the stratified 5-fold cross validation technique for both model building and validation
- Utilize effective standardization techniques and hyperparameter tuning
- Incorporate sklearn-pipelines to streamline the workflow
- Apply appropriate regularization techniques to address the bias-variance trade-off
- Create five folds in the data, and introduce a variable to identify the folds
- Develop Gradient Boost model and determine the variable importance scores, and identify the redundant variables

```
# lets first seperate input and output data
data = Customer_details_noQ[['children', 'charges', 'bmi', 'hbalc',
                             'heart_issues', 'any_transplants', 'cancer_history',
                             'numberofmajorsurgeries', 'smoker', 'city_tier_ord',
                             'hospital_tier_ord', 'State_ID_R1013', 'State_ID_R1011',
                             'State_ID_R1012', 'age', 'gender']]

final_data = pd.get_dummies(data, drop_first=True, dtype='int')

X = final_data.drop(['charges'], axis=1)
y = final_data[['charges']]

X.head()
```

	children	bmi	hbalc	numberofmajorsurgeries	city_tier_ord	\
0	0	17.58	4.51	1.0	2	
1	0	17.60	4.39	1.0	0	
2	0	16.47	6.35	1.0	0	
3	0	17.70	6.28	1.0	2	
4	0	22.34	5.57	1.0	2	

	hospital_tier_ord	State_ID_R1013	State_ID_R1011	State_ID_R1012
age \				
0	1	1	0	0
32				
1	1	1	0	0
32				
2	1	1	0	0
31				
3	2	1	0	0
32				
4	2	1	0	0
26				

	heart_issues_yes	any_transplants_yes	cancer_history_Yes
smoker_yes \			

0	0	0	0
0			
1	0	0	0
0			
2	0	0	1
0			
3	0	0	0
0			
4	0	0	0
0			

	gender_Male
0	1
1	1
2	0
3	1
4	1

setting up a pipe line

```
from sklearn.linear_model import SGDRegressor, Ridge
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import
KFold, StratifiedKFold, RandomizedSearchCV, GridSearchCV
from sklearn.metrics import mean_squared_error
```

```
pipeline = Pipeline(steps=[('scaler', StandardScaler()),
                           ('regression', Ridge())])
```

Defining the parameter for hyper parameter tuning

```
parameters = {'regression__alpha' : [0.001, 0.01, 0.1, 1, 10, 100]}
```

creating k-fold objects

```
kfold = KFold(n_splits=5, shuffle=True, random_state=3)
```

creating the gradient search objects

```
model_ridge =
GridSearchCV(pipeline, param_grid=parameters, cv=kfold, scoring='neg_mean
_squared_error')
```

```
model_ridge.fit(X, y)
```

```
GridSearchCV(cv=KFold(n_splits=5, random_state=3, shuffle=True),
             estimator=Pipeline(steps=[('scaler', StandardScaler()),
                                       ('regression', Ridge())]),
             param_grid={'regression__alpha': [0.001, 0.01, 0.1, 1,
10, 100]},
             scoring='neg_mean_squared_error')
```

```
model_ridge.best_params_
```

```
{'regression__alpha': 10}
model_ridge.best_estimator_
Pipeline(steps=[('scaler', StandardScaler()), ('regression',
Ridge(alpha=10))])
```

Gradient Boosting Algorithms

```
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.model_selection import train_test_split

X_train,X_test,y_train,y_test = train_test_split(X,y)

model = GradientBoostingRegressor()
model.fit(X_train,y_train)

GradientBoostingRegressor()

model.feature_importances_

array([4.38853904e-03, 1.14916455e-01, 3.89144324e-03, 6.40714004e-04,
       1.54577429e-04, 2.03448674e-02, 4.78820728e-03, 8.91851976e-03,
       3.68398264e-04, 9.87755494e-02, 9.63030890e-05, 0.00000000e+00,
       8.27881010e-05, 7.41043275e-01, 1.59036286e-03])
```

Variable importance

```
pd.DataFrame({'features' : model.feature_names_in_, 'importance' :
model.feature_importances_}).sort_values('importance',ascending=False)
```

	features	importance
13	smoker_yes	0.741043
1	bmi	0.114916
9	age	0.098776
5	hospital_tier_ord	0.020345
7	State_ID_R1011	0.008919
6	State_ID_R1013	0.004788
0	children	0.004389
2	hbalc	0.003891
14	gender_Male	0.001590
3	numberofmajorsurgeries	0.000641
8	State_ID_R1012	0.000368
4	city_tier_ord	0.000155
10	heart_issues_yes	0.000096
12	cancer_history_Yes	0.000083
11	any_transplants_yes	0.000000

```
# train_score
model.score(X_train,y_train)

0.9355591853437756

# test score
model.score(X_test,y_test)

0.9132710744585373
```

Problem - 3

Estimate the cost of hospitalization for Christopher, Ms. Jayna (Date of birth 12/28/1988; height 170 cm; and weight 85 kgs). She lives with her partner and two children in a tier-1 city, and her state's State ID is R1011. She was found to be nondiabetic (HbA1c = 5.8). She smokes but is otherwise healthy. She has had no transplants or major surgeries. Her father died of lung cancer. Hospitalization costs will be estimated using tier-1 hospitals.

```
Customer_details_noQ.columns

Index(['customer_id', 'year', 'month', 'date', 'children', 'charges',
      'hospital_tier', 'city_tier', 'state_id', 'bmi', 'hbalc',
      'heart_issues', 'any_transplants', 'cancer_history',
      'numberofmajorsurgeries', 'smoker', 'name',
      'hospital_tier_ord',
      'city_tier_ord', 'State_ID_R1013', 'State_ID_R1011',
      'State_ID_R1012',
      'age', 'title', 'gender'],
      dtype='object')

pred_data = pd.DataFrame({
    'name': ['Christopher, Ms. Jayna'],
    'dob': ['12/28/1988'],
    'children': [2],
    'bmi': [85/(1.7**2)],
    'hbalc': [5.8],
    'numberofmajorsurgeries': [0],
    'city_tier_ord': [1],
    'hospital_tier_ord': [1],
    'state_id_R1013': [0],
    'state_id_R1011': [1],
    'state_id_R1012': [0],
    'age': [36],
    'heart_issues_yes': [0],
    'any_transplants_yes': [0],
    'cancer_history_Yes': [1],
    'smoker_yes': [1],
```

```

    'gender_male':[0]
}))
X.columns
Index(['children', 'bmi', 'hbalc', 'numberofmajorsurgeries',
      'city_tier_ord',
      'hospital_tier_ord', 'State_ID_R1013', 'State_ID_R1011',
      'State_ID_R1012', 'age', 'heart_issues_yes',
      'any_transplants_yes',
      'cancer_history_Yes', 'smoker_yes', 'gender_Male'],
      dtype='object')

pred_data

```

	name	dob	children	bmi	hbalc	\
0	Christopher, Ms. Jayna	12/28/1988	2	29.411765	5.8	

	numberofmajorsurgeries	city_tier_ord	hospital_tier_ord
state_id_R1013 \			
0	0	1	1
0			

	state_id_R1011	state_id_R1012	age	heart_issues_yes
any_transplants_yes \				
0	1	0	36	0
0				

	cancer_history_Yes	smoker_yes	gender_male
0	1	1	0

```

pred_data['dob'] = pd.to_datetime(pred_data.dob,errors='coerce')
age=2024 - pred_data.dob.dt.year
age
0    36
Name: dob, dtype: int64

test_data =pred_data.drop(['name','dob'],axis=1)

```

predicting the charges of test_data using best model

```

model.predict(test_data)
array([27777.69719166])

```

----- End of Machine Learning Part -----